

Connecting to an RDS instance

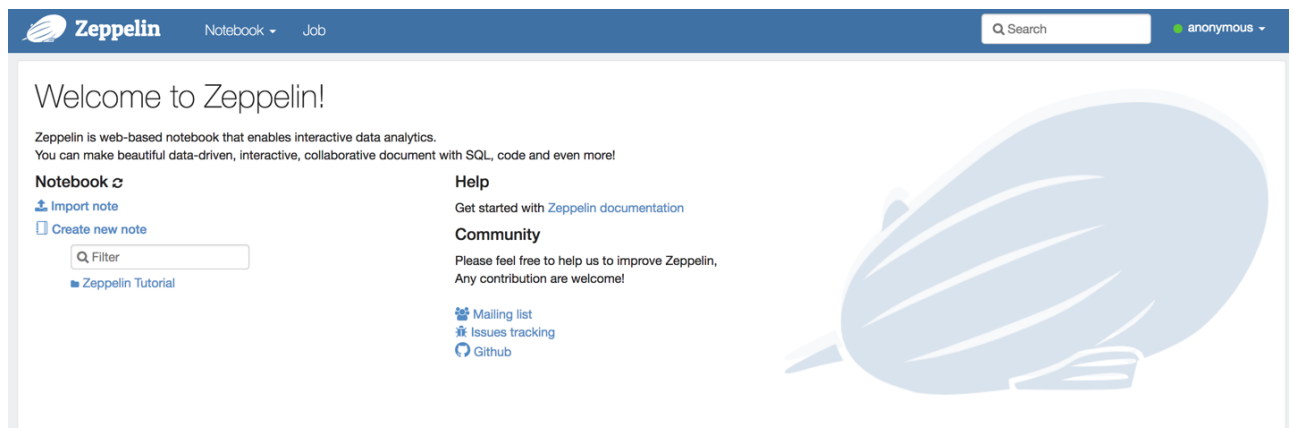
This guide will setup connections to an AWS RDS instance that can be used inside a Docker-run Zeppelin image.

Prerequisites

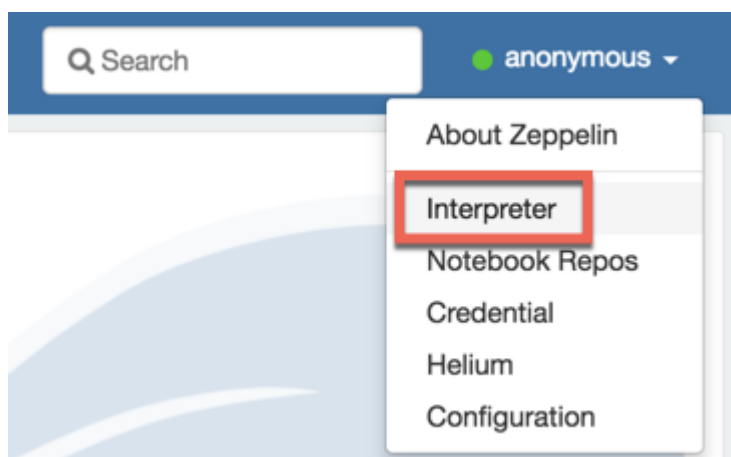
- Install Docker on your operating system. Links and instructions can be found in the [README.md](#)
- Create an RDS instance in AWS.
- Update RDS security data to accept all traffic, both Inbound and Outbound.
- Create a table with data inserted. This will allow you to test the connection.

Instructions

- Run the docker image `docker run -p 8080:8080 --rm --name zeppelin trilogyed/zepl:1.0.0`.
- **Note:** The image will be downloaded the first time it is run, which will take a considerable amount of time. After the image has been downloaded, the container will start up much faster in subsequent runs.
- Navigate to `localhost:8080`.



- Click **Profile** on top right and select **Interpreter**.



- Click **+Create**.

- Enter an interpreter name and select **jdbc** from the "Interpreter group".
- Enter the password for the RDS in **default.password**.
- Add the AWS endpoint to **default.url** before the port and the database name after port.
- Enter the username for the RDS instance in **default.user**.

Properties		
name	value	action
common.max_count	<input type="text" value="1000"/>	<input type="button" value="✕"/>
default.completer.schemaFilters	<input type="text"/>	<input type="button" value="✕"/>
default.completer.ttlInSeconds	<input type="text" value="120"/>	<input type="button" value="✕"/>
default.driver	<input type="text" value="org.postgresql.Driver"/>	<input type="button" value="✕"/>
default.password	<input type="password" value="....."/>	<input type="button" value="✕"/>
default.precode	<input type="text"/>	<input type="button" value="✕"/>
default.splitQueries	<input type="checkbox"/>	<input type="button" value="✕"/>
default.statementPrecode		<input type="button" value="✕"/>
default.url	<input type="text" value="jdbc:postgresql://<aws endpoint>:5432/<db name>"/>	<input type="button" value="✕"/>
default.user	<input type="text" value="<db user>"/>	<input type="button" value="✕"/>

- Scroll down and click **Save**.
- Navigate back to the main page and click **Create new note**.
- Enter a name for the note, and then select the name of the interpreter just created as the default interpreter. Click **Create**.

Create New Note

Note Name

Default Interpreter

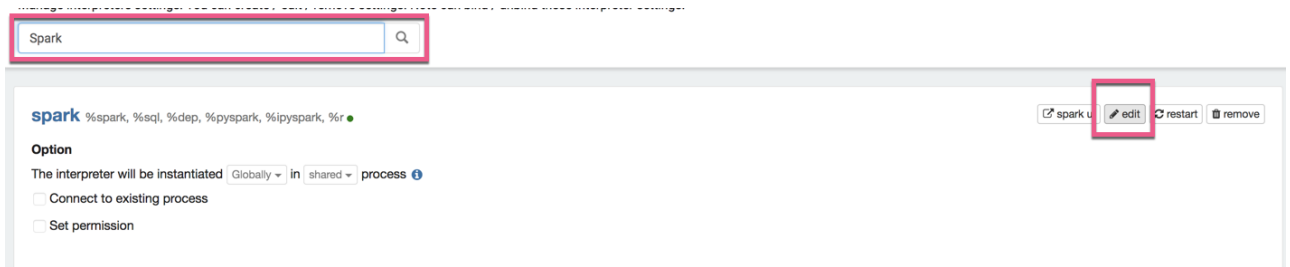
Use '/' to create folders. Example: /NoteDirA/Note1

Create

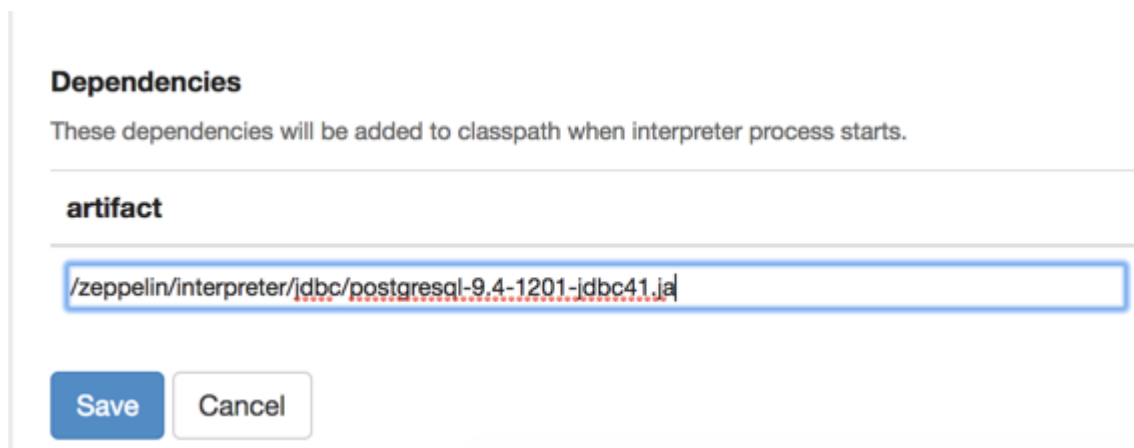
- Run a **SELECT * FROM <table name>** from a cell to confirm the connection works.

Adding a Postgres Drive to ZEPL

- On the interpreters screen, search for **Spark**.



- Click **Edit** and scroll down to Dependencies.
- Under "artifact" enter the path `/zeppelin/interpreter/jdbc/postgresql-9.4-1201-jdbc41.jar` then click **Save**.



- Spark will now be able to connect to an RDS instance using a JDBC URL.
- Test that the connection works by running the following code, using a table currently available in your RDS:

```
%pyspark
jdbc_url = "jdbc:postgresql://<rds endpoint>:5432/<DB name>"
config = {"user": "root", "password": "<password>"}
my_df = spark.read.jdbc(url=jdbc_url, table='<table name>',
properties=config)
my_df.limit(10).show()
```